

TextonBoost : Joint Appearance, Shape
and Context Modeling for Multi-class Object
Recognition and Segmentation

Jamie Shotton², JoHNWinn¹, Carsten Rother¹, and Antonio Criminisi¹
¹Microsoft Research Ltd., Cambridge, UK
{jwinn, carrot, antcrim}@microsoft.com
²Department of Engineering,
University of Cambridge
jdjs2@cam.ac.uk

Abstract. This paper proposes a new approach to learning a discriminative model of object classes, incorporating appearance, shape and context information efficiently. The learned model is used for automatic visual recognition and semantic segmentation of photographs. Our discriminative model exploits novel features, based on textons, which jointly model shape and texture. Unary classification and feature selection is achieved using shared boosting to give an efficient classifier which can be applied to a large number of classes. Accurate image segmentation is achieved by incorporating these classifiers in a conditional random field. Efficient training of the model on very large datasets is achieved by exploiting both random feature selection and piecewise training methods. High classification and segmentation accuracy are demonstrated on three different databases: i) our own 21-object class database of photographs of real objects viewed under general lighting conditions, poses and viewpoints, ii) the 7-class Corel subset and iii) the 7-class Sowerby database used in [1]. The proposed algorithm gives competitive results both for highly textured (e.g. grass, trees), highly structured (e.g. cars, faces, bikes, aeroplanes) and articulated objects (e.g. body, cow).

1

Weakly Supervised Learning of Part-Based
Spatial Models for Visual Object Recognition
David J. Crandall and Daniel P. Huttenlocher
Cornell University, Ithaca, NY 14850, USA
{crandall, dph}@cs.cornell.edu

Abstract. In this paper we investigate a new method of learning part-based models for visual object recognition, from training data that only provides information about class membership (and not object location or configuration). This method learns both a model of local part appearance and a model of the spatial relations between those parts. In contrast, other work using such a weakly supervised learning paradigm has not considered the problem of simultaneously learning appearance and spatial models. Some of these methods use a "bag" model where only part appearance is considered whereas other methods learn spatial models but only given the output of a particular feature detector. Previous techniques for learning both part appearance and spatial relations have instead used a highly supervised learning process that provides substantial information about object part location. We show that our weakly supervised technique produces better results than these previous highly supervised methods. Moreover, we investigate the degree to which both richer spatial models and richer appearance models are helpful in improving recognition performance. Our results show that while both spatial and appearance information can be useful, the effect on performance depends substantially on the particular object class and on the difficulty of the test dataset.

1

Hyperfeatures - Multilevel Local Coding
for Visual Recognition
Ankur Agarwal and Bill Triggs

GRA VIR-INRIA-CNRS, 655 Avenue de l'Europe, Montbonnot 38330, France

{Ankur.Agarwal, Bill.Triggs }@inrialpes.fr
<http://www.inrialpes.fr/lear/people/{agarwal, triggs }>
 Abstract. Histograms of local appearance descriptors are a popular representation for visual recognition. They are highly discriminant and have good resistance to local occlusions and to geometric and photometric variations, but they are not able to exploit spatial co-occurrence statistics at scales larger than their local input patches. We present a new multilevel visual representation, 'hyperfeatures', that is designed to remedy this. The starting point is the familiar notion that to detect object parts, in practice it often suffices to detect co-occurrences of more local object fragments - a process that can be formalized as comparison (e.g. vector quantization) of image patches against a codebook of known fragments, followed by local aggregation of the resulting codebook membership vectors to detect co-occurrences. This process converts local collections of image descriptor vectors into somewhat less local histogram vectors - higher-level but spatially coarser descriptors. We observe that as the output is again a local descriptor vector, the process can be iterated, and that doing so captures and codes ever larger assemblies of object parts and increasingly abstract or 'semantic' image properties. We formulate the hyperfeatures model and study its performance under several different image coding methods including clustering based Vector Quantization, Gaussian Mixtures, and combinations of these with Latent Dirichlet Allocation. We find that the resulting high-level features provide improved performance in several object image and texture image classification tasks.

1

 Riemannian Manifold Learning for Nonlinear
 Dimensionality Reduction
 Tony Lin^{1,2}, Hongbin Zha¹, and Sang Uk Lee²
¹National Laboratory on Machine Perception,
 Peking University, Beijing 100871, China
 {lintong, zha}@cis.pku.edu.cn
²School of Electrical Engineering,
 Seoul National University, Seoul 151-742, Korea
 sanguk@ipl.snu.ac.kr

Abstract. In recent years, nonlinear dimensionality reduction (NLDR) techniques have attracted much attention in visual perception and many other areas of science. We propose an efficient algorithm called Riemannian manifold learning (RML). A Riemannian manifold can be constructed in the form of a simplicial complex, and thus its intrinsic dimension can be reliably estimated. Then the NLDR problem is solved by constructing Riemannian normal coordinates (RNC). Experimental results demonstrate that our algorithm can learn the data's intrinsic geometric structure, yielding uniformly distributed and well organized low-dimensional embedding data.

1

 Controlling Sparseness in
 Non-negative Tensor Factorization
 Matthias Heiler and Christoph Schnörr
 Computer Vision, Graphics, and Pattern Recognition Group,
 Department of Mathematics and Computer Science,
 University of Mannheim, 68131 Mannheim, Germany
 {heiler, schnoerr}@uni-mannheim.de
 Abstract. Non-negative tensor factorization (NTF) has recently been

proposed as sparse and efficient image representation (Welling and Weber, *Patt. Rec. Let.*, 2001). Until now, sparsity of the tensor factorization has been empirically observed in many cases, but there was no systematic way to control it. In this work, we show that a sparsity measure recently proposed for non-negative matrix factorization (Hoyer, *J. Mach. Learn. Res.*, 2004) applies to NTF and allows precise control over sparseness of the resulting factorization. We devise an algorithm based on sequential conic programming and show improved performance over classical NTF codes on artificial and on real-world data sets.

1

Conditional Infomax Learning: An Integrated
Framework for Feature Extraction and Fusion

Dahua Lin¹, Linde Xie¹, and Tao Xiang^{1,2}

¹Dept. of Information Engineering,

The Chinese University of Hong Kong, Hong Kong, China

dhlin4@ie.cuhk.edu.hk

²Microsoft Research Asia, Beijing, China

xitang@microsoft.com

Abstract. The paper introduces a new framework for feature learning in classification motivated by information theory. We first systematically study the information structure and present a novel perspective revealing the two key factors in information utilization: class-relevance and redundancy. We derive a new information decomposition model where a novel concept called class-relevant redundancy is introduced. Subsequently a new algorithm called Conditional Informative Feature Extraction is formulated, which maximizes the joint class-relevant information by explicitly reducing the class-relevant redundancies among features. To address the computational difficulties in information-based optimization, we incorporate Parzen window estimation into the discrete approximation of the objective function and propose a Local Active Region method which substantially increases the optimization efficiency. To effectively utilize the extracted feature set, we propose a Bayesian MAP formulation for feature fusion, which unifies Laplacian Sparse Prior and Multivariate Logistic Regression to learn a fusion rule with good generalization capability. Realizing the inefficiency caused by separate treatment of the extraction stage and the fusion stage, we further develop an improved design of the framework to coordinate the two stages by introducing a feedback from the fusion stage to the extraction stage, which significantly enhances the learning efficiency. The results of the comparative experiments show remarkable improvements achieved by our framework.

1

Degen Generalized Cylinders and Their Properties

Liangliang Cao¹, Jianzhuang Liu¹, and Xiaoou Tang^{1,2}

¹Department of Information Engineering, The Chinese University of Hong Kong, Hong Kong, China

{llcao, jzliu, xtang}@ie.cuhk.edu.hk

²Microsoft Research Asia, Beijing, China

xitang@microsoft.com

Abstract. Generalized cylinder (GC) has played an important role in computer vision since it was introduced in the 1970s. While studying GC models in human visual perception of shapes from contours, Marr assumed that GC's limbs are planar curves. Later, Koenderink and Ponce pointed out that this assumption does not hold in general by giving some examples. In this paper, we show that straight homogeneous generalized cylinders (SHGCs) and tori (a kind of curved GCs) have planar limbs when viewed from points on specific straight lines. This property leads us to the definition and investigation of a new class of GCs, with the help of the surface model proposed by Degen for geometric modeling. We call them Degen generalized cylinders (DGCs), which include SHGCs, tori, quadrics, cyc

lides, and more other GCs into one model. Our rigorous discussion is based on projective geometry and homogeneous coordinates. We present some invariant properties of DGCs that reveal the relations among the planar limbs, axes, and contours of DGCs. These properties are useful for recovering DGC descriptions from image contours as well as for some other tasks in computer vision.

1

Geodesics Between 3D Closed Curves Using Path-Straightening

Eric Klassen¹ and Anuj Srivastava²

¹Department of Mathematics, Florida State University, Tallahassee, FL 32306

²Department of Statistics, Florida State University, Tallahassee, FL 32306

Abstract. In order to analyze shapes of continuous curves in \mathbb{R}^3 , we parameterize them by arc-length and represent them as curves on a unit two-sphere. We identify the subset denoting the closed curves, and study its differential geometry. To compute geodesics between any two such curves, we connect them with an arbitrary path, and then iteratively straighten this path using the gradient of an energy associated with this path. The limiting path of this path-straightening approach is a geodesic. Next, we consider the shape space of these curves by removing shape-preserving transformations such as rotation and re-parametrization. To construct a geodesic in this shape space, we construct the shortest geodesic between the all possible transformations of the two end shapes; this is accomplished using an iterative procedure. We provide step-by-step descriptions of all the procedures, and demonstrate them with simple examples.

1

Robust Homography Estimation from Planar

Contours Based on Convexity

Alberto Ruiz¹, Pedro E. L´opez de Teruel², and Lorenzo Fern´andez²

¹Dept. Inform´atica y Sistemas, University of Murcia

²Dept. Tecnolog´ıa e Ingenier´ıa de Computadores, University of Murcia

aruiz@um.es, {pedroel, lfmaimo}@ditec.um.es

Abstract. We propose a homography estimation method from the contours of planar regions. Standard projective invariants such as cross ratios or canonical frames based on hot points obtained from local differential properties are extremely unstable in real images suffering from pixelization, thresholding artifacts, and other noise sources. We explore alternative constructions based on global convexity properties of the contour such as discrete tangents and concavities. We show that a projective frame can be robustly extracted from arbitrary shapes with at least one appreciable concavity. Algorithmic complexity and stability are theoretically discussed and experimentally evaluated in a number of real applications including projective shape matching, alignment and pose estimation. We conclude that the procedure is computationally efficient and notably robust given the ill-conditioned nature of the problem.

1

Detecting Instances of Shape Classes That Exhibit Variable Structure

Vassilis Athitsos¹, Jingbin Wang², Stan Sclaroff², and Margaret B. Eick²

¹Siemens Corporate Research, Princeton, NJ 08540, USA

²Computer Science Department, Boston University, Boston, MA 02215, USA

Abstract. This paper proposes a method for detecting shapes of variable structure in images with clutter. The term “variable structure” means that some shape parts can be repeated an arbitrary number of times, some parts can be optional, and some parts can have several alternative appearances. The particular variation

of the shape structure that occurs in a given image is not known a priori. Ex-is

ting computer vision methods, including deformable model methods, were not designed to detect shapes of variable structure; they may only be used to detect shapes that can be decomposed into a fixed, a priori known, number of parts. The proposed method can handle both variations in shape structure and variations in the appearance of individual shape parts. A new class of shape models is introduced, called Hidden State Shape Models, that can naturally represent shapes of variable structure. A detection algorithm is described that finds instances of such

shapes in images with large amounts of clutter by finding globally optimal correspondences between image features and shape models. Experiments with real images demonstrate that our method can localize plant branches that consist of an a priori unknown number of leaves and can detect hands more accurately than a hand detector based on the chamfer distance.

1

Direct Solutions for Computing Cylinders from Minimal Sets of 3D Points

Christian Beder and Wolfgang Förstner

Institute for Photogrammetry,
Bonn University, Germany

{beder, wf}@ipb.uni-bonn.de

Abstract. Efficient direct solutions for the determination of a cylinder from points are presented. The solutions range from the well known direct solution of a quadric to the minimal solution of a cylinder with five points. In contrast to the approach of G. Roth and M. D. Levine (1990), who used polynomial bases for representing the geometric entities, we use algebraic constraints on the quadric representing the cylinder. These solutions for six to eight points directly determine all the cylinder parameters in one step: (1) The eight-point-solution, similar to the estimation of the fundamental matrix, requires to solve for the roots of a 3rd-order polynomial. (2) The seven-point-solution, similar to the six-point-solution for the relative orientation by J. Philip (1996), yields a linear equation system. (3) The six-point-solution, similar to the five-point-solution for the relative orientation by D. Nister (2003), yields a ten-by-ten eigenvalue problem. The new minimal five-point-solution first determines the direction and then the position and the radius of the cylinder. The search for the zeros of the resulting 6th order polynomials is efficiently realized using 2D-Bernstein polynomials. Also direct solutions for the special cases with the axes of the cylinder parallel to a coordinate plane or axis are given. The method is used to find cylinders in range data of an industrial site.

1

Estimation of Multiple Periodic Motions from Video

Alexia Briassouli and Narendra Ahuja

Beckman Institute, University of Illinois, Urbana-Champaign,
405 N Matthews, Urbana, IL, 61801

{briassou, ahuja}@vision.ai.uiuc.edu

Abstract. The analysis of periodic or repetitive motions is useful in many applications, both in the natural and the man-made world. An important example is the recognition of human and animal activities. Existing methods for the analysis of periodic motions first extract motion trajectories, e.g. via correlation, or feature point matching. We present a new approach, which takes advantage of both the frequency and spatial information of the video. The 2D spatial Fourier transform is applied to each frame, and time-frequency distributions are then used to estimate the time-varying object motions. Thus, multiple periodic trajectories are extracted and their periods are estimated. The period information is finally used to segment the periodically moving objects. Unlike existing

methods, our approach estimates multiple periodicities simultaneously, it is robust to deviations from strictly periodic motion, and estimates periodicities superposed on translations. Experiments with synthetic and real sequences display the capabilities and limitations of this approach. Supplementary material is provided, showing the video sequences used in the experiments.

1

Robust Multi-body Motion Tracking Using Commute Time Clustering

Huaijun Qiu and Edwin R. Hancock

Department of Computer Science, University of York
York, YO10 5DD, UK

Abstract. The presence of noise renders the classical factorization method almost impractical for real-world multi-body motion tracking problems. The main problem stems from the effect of noise on the shape interaction matrix, which loses its block-diagonal structure and as a result the assignment of elements to objects becomes difficult. The aim in this paper is to overcome this problem using graph-spectral embedding and the k-means algorithm. To this end we develop a representation based on the commute time between nodes on a graph. The commute time (i.e. the expected time taken for a random walk to travel between two nodes and return) can be computed from the Laplacian spectrum using the discrete Green's function, and is an important property of the random walk on a graph. The commute time is a more robust measure of the proximity of data than the raw proximity matrix. Our embedding procedure preserves commute time, and is closely akin to kernel PCA, the Laplacian eigenmap and the diffusion map. We illustrate the results both on the synthetic image sequences and real-world video sequences, and compare our results with several alternative methods.

1

A Tuned Eigenspace Technique for Articulated Motion Recognition

M. Masudur Rahman and Antonio Robles-Kelly

National ICT Australia, RSISE Bldg. 115, ANU, ACT 0200, Australia

Masud.Rahman@ {rsise.anu.edu.au, nicta.com.au }

Antonio.Robles-Kelly@ {anu.edu.au, nicta.com.au }

Abstract. In this paper, we introduce a tuned eigenspace technique so as to classify human motion. The method presented here overcomes those problems related to articulated motion and dress texture effects by learning various human motions in terms of their sequential postures in an eigenspace. In order to cope with the variability inherent to articulated motion, we propose a method to tune the set of sequential eigenspaces. Once the learnt tuned eigenspaces are at hand, the recognition task then becomes a nearest-neighbor search over the eigenspaces. We show how our tuned eigenspace method can be used for purposes of real-world and synthetic pose recognition. We also discuss and overcome the problem related to clothing texture that occurs in real-world data, and propose a background subtraction method to employ the method in out-door environment. We provide results on synthetic imagery for a number of human poses and illustrate the utility of the method for the purposes of human motion recognition.

1

1

Real-Time Non-rigid Shape Recovery Via Active Appearance Models for Augmented Reality

Jianke Zhu, Steven C.H. Hoi, and Michael R. Lyu

Department of Computer Science & Engineering,
Chinese University of Hong Kong,

Shatin, Hong Kong

{jkzhu, chhoi, lyu }@cse.cuhk.edu.hk

Abstract. One main challenge in Augmented Reality (AR) applications is to keep track of video objects with their movement, orientation, size, and position accurately. This poses a challenging task to recover non-rigid shape and global pose in real-time AR applications. This paper proposes a novel two-stage scheme for online non-rigid shape recovery toward AR applications using Active Appearance Models (AAMs). First, we construct 3D shape models from AAMs offline, which do not involve processing of the 3D scan data. Based on the computed 3D shape models, we propose an efficient online algorithm to estimate both 3D pose and non-rigid shape parameters via local bundle adjustment for building up point correspondences. Our approach, without manual intervention, can recover the 3D non-rigid shape effectively from either real-time video sequences or single image. The recovered 3D pose parameters can be used for AR registrations. Furthermore, the facial feature can be tracked simultaneously, which is critical for many face related applications. We evaluate our algorithms on several video sequences. Promising experimental results demonstrate our proposed scheme is effective and significant for real-time AR applications.

1

A Fluid Motion Estimator for Schlieren

Image Velocimetry

Elise Arnaud¹, Etienne M´emin², Robert O’Sosa³, and Guillermo Artana³

¹Disi, Universit`a di Genova, 16146 Genova, Italy

arnaud@disi.unige.it

²IRISA, Universit´e de Rennes 1, 35 042 Rennes Cedex, France

memin@irisa.fr

³Facultad de Ingenier´ıa, Universidad de Buenos Aires

Buenos Aires 1412, Argentina

{rsosa, gartana }@fi.uba.ar

Abstract. In this paper, we address the problem of estimating the motion of fluid flows that are visualized through a Schlieren system. Such a system is well known in fluid mechanics as it enables the visualization of unseeded flows. As the resulting images exhibit very low photometric contrasts, classical motion estimation methods based on the brightness consistency assumption (correlation-based approaches, optical flow methods) are completely inefficient. This work aims at proposing a sound energy based estimator dedicated to these particular images. The energy function to be minimized is composed of (a) a novel data term describing the fact that the observed luminance is linked to the gradient of the fluid density and (b) a specific div curl regularization term. The relevance of our estimator is demonstrated on real-world sequences.

1

Bilateral Filtering-Based Optical Flow Estimation

with Occlusion Detection

Jiangjian Xiao, Hui Cheng, Harpreet Sawhney,

Cen Rao, and Michael Isnardi

Sarnoff Corporation

{jxiao, hcheng, hsawhney, crao, misnardi }@sarnoff.com

Abstract. Using the variational approaches to estimate optical flow between two frames, the flow discontinuities between different motion fields are usually not distinguished even when an anisotropic diffusion operator is applied. In this pa-

-

per, we propose a multi-cue driven adaptive bilateral filter to regularize the flow computation, which is able to achieve the smoothly varied optical flow field with highly desirable motion discontinuities. First, we separate the traditional one-s

w

tep

variational updating model into a two-step ■ltering-based updating model. Then, employing our occlusion detector, we reformulate the energy functional of optical ■low estimation by explicitly introducing an occlusion term to balance the energy loss due to the occlusion or mismatches. Furthermore, based on the two-step updating framework, a novel multi-cue driven bilateral ■lter is proposed to substitute the original anisotropic diffusion process, and it is able to adaptively

control the diffusion process according to the occlusion detection, image intensity dissimilarity, and motion dissimilarity. After applying our approach on various video sources (movie and TV) in the presence of occlusion, motion blurring

, non-rigid deformation, and weak texture, we generate a spatial-coherent ■low ■eld between each pair of input frames and detect more accurate ■low discontinuities along the motion boundaries.

1

Geometry and Kinematics with Uncertain Data

Christian Perwass, Christian Gebken, and Gerald Sommer

Institut f"ur Informatik, CAU Kiel,

Christian-Albrechts-Platz 4, 24118 Kiel, Germany

{chp, chg, gs}@ks.informatik.uni-kiel.de

Abstract. In Computer Vision applications, one usually has to work with uncertain data. It is therefore important to be able to deal with uncertain geometry and uncertain transformations in a uniform way. The Geometric Algebra of conformal space offers a unifying framework to treat not only geometric entities like points, lines, planes, circles and spheres, but also transformations like reflection, inversion, rotation and translation. In this text we show how the un-

certainty of all elements of the Geometric Algebra of conformal space can be appropriately described by covariance matrices. In particular, it will be shown that it is advantageous to represent uncertain transformations in Geometric Alge-

bra as compared to matrices. Other important results are a novel pose estimation approach, a uniform framework for geometric entity ■tting and triangulation, the testing of uncertain tangency relations and the treatment of catadioptric cameras with parabolic mirrors within this framework. This extends previous work by Forstner and Heuel from points, lines and planes to non-linear geometric entities

and transformations, while keeping the linearity of the estimation method. We give a theoretical description of our approach and show exemplary applications.

1

Euclidean Structure from $N \geq 2$ Parallel Circles:

Theory and Algorithms

Pierre Gurdjos¹, Peter Sturm², and Yihong Wu³

¹IRIT-TCI, UPS, 118 route de Narbonne,

31062 Toulouse, cedex 9, France

Pierre.Gurdjos@irit.fr

²PERCEPTION, INRIA Rhône-Alpes,

655, avenue de l'Europe, 38330 Montbonnot, France

Peter.Sturm@inrialpes.fr

³NLPR-IA, Chinese Academy of Sciences, P.O. Box 2728,

No. 95 East Road of Zhong Guan Cun, Beijing 100080, China

yhwu@nlpr.ia.ac.cn

Abstract. Our problem is that of recovering, in one view, the 2D Euclidean structure, induced by the projections of N parallel circles. This structure is a prerequisite for camera calibration and pose computation. Until now, no general method has been described for $N > 2$. The main

contribution of this work is to state the problem in terms of a system of linear equations to solve. We give a closed-form solution as well as bundle adjustment-like refinements, increasing the technical applicability and numerical stability. Our theoretical approach generalizes and extends all those described in existing works for $N=2$ in several respects, as we can treat simultaneously pairs of orthogonal lines and pairs of circles within a unified framework. The proposed algorithm may be easily implemented, using well-known numerical algorithms. Its performance is illustrated by simulations and experiments with real images.

1

Overconstrained Linear Estimation of Radial Distortion and Multi-view Geometry

R. Matt Steele and Christopher Jaynes

Center for Visualization and Virtual Environments,
University of Kentucky, Lexington, Kentucky, USA

Abstract. This paper introduces a new method for simultaneous estimation of lens distortion and multi-view geometry using only point correspondences. The new technique has significant advantages over the current state-of-the art in that it makes more effective use of correspondences arising from any number of views. Multi-view geometry in the presence of lens distortion can be expressed as a set of point correspondence constraints that are quadratic in the unknown distortion parameter. Previous work has demonstrated how the system can be solved efficiently as a quadratic eigenvalue problem by operating on the normal equations of the system. Although this approach is appropriate for situations in which only a minimal set of matchpoints are available, it does not take full advantage of extra correspondences in overconstrained situations, resulting in significant bias and many potential solutions. The new technique directly operates on the initial constraint equations and solves the quadratic eigenvalue problem in the case of rectangular matrices. The method is shown to contain significantly less bias on both controlled and real-world data and, in the case of a moving camera where additional views serve to constrain the number of solutions, an accurate estimate of both geometry and distortion is achieved.

1

Camera Calibration with Two Arbitrary Coaxial Circles

Carlo Colombo, Dario Comanducci, and Alberto Del Bimbo

Dipartimento di Sistemi e Informatica,
Via S. Marta 3, 50139 Firenze, Italy

{colombo, comandu, delbimbo}@dsi.unifi.it

Abstract. We present an approach for camera calibration from the image of at least two circles arranged in a coaxial way. Such a geometric configuration arises in static scenes of objects with rotational symmetry or in scenes including generic objects undergoing rotational motion around a fixed axis. The approach is based on the automatic localization of a surface of revolution (SOR) in the image, and its use as a calibration artifact. The SOR can either be a real object in a static scene, or a "virtual surface" obtained by frame superposition in a rotational sequence. This provides a unified framework for calibration from single images of SORs or from turntable sequences. Both the internal and external calibration parameters (square pixels model) are obtained from two or more imaged cross sections of the SOR, whose apparent contour is also exploited to obtain a better calibration accuracy. Experimental results show that this calibration approach is accurate enough for several vision applications, encompassing 3D realistic model acquisition from single images, and desktop 3D object scanning.

1

Molding Face Shapes by Example
Ira Kemelmacher and Ronen Basri
Dept. of Computer Science and Applied Math.,
The Weizmann Institute of Science,
Rehovot 76100, Israel

Abstract. Human faces are remarkably similar in global properties, including size, aspect ratios, and locations of main features, but can vary considerably in details across individuals, gender, race, or due to facial expression. We propose a novel method for 3D shape recovery of a face from a single image using a single 3D reference model of a different person's face. The method uses the input image as a guide to mold the reference model to reach a desired reconstruction. Assuming Lambertian reflectance and rough alignment of the input image and reference model, we seek shape, albedo, and lighting that best fit the image while preserving the rough structure of the model. We demonstrate our method by providing accurate reconstructions of novel faces overcoming significant differences in shape due to gender, race, and facial expressions.

1

Reconstruction of Canal Surfaces from Single
Images Under Exact Perspective
Vincenzo Caglioti and Alessandro Giusti
Politecnico di Milano
{caglioti, giusti}@elet.polimi.it

Abstract. This paper addresses the reconstruction of canal surfaces from single images. A canal surface is obtained as the envelope of a family of spheres of constant radius, whose center is swept along a spacecurve, called axis. Previous studies either used approximate relationships (quasi-invariants), or they addressed the recognition based on a geometric model. In this paper we show that, under broad conditions, canal surfaces can be reconstructed from single images under exact perspective. In particular, canal surfaces with planar axis can even be reconstructed from a single fully-uncalibrated image. An automatic reconstruction method has been implemented. Simulations and experimental results on real images are also presented.

1

Subspace Estimation Using Projection Based
M-Estimators over Grassmann Manifolds
Raghav Subbarao and Peter Meer
Department of Electrical and Computer Engineering,
Rutgers University, Piscataway, NJ 08854, USA
{rsubbara, meer}@caip.rutgers.edu

Abstract. We propose a solution to the problem of robust subspace estimation using the projection based M-estimator. The new method handles more outliers than inliers, does not require a user defined scale of the noise affecting the inliers, handles noncentered data and nonorthogonal subspaces. Other robust methods like RANSAC, use an input for the scale, while methods for subspace segmentation, like GPCA, are not robust. Synthetic data and three real cases of multibody factorization show the superiority of our method, in spite of user independence.

1

Anastassia Angelopoulou¹, José García-Rodríguez², and Alexandra Psarrou¹
¹Harrow School of Computer Science, University of Westminster,
Harrow HA1 3TP, United Kingdom
²Departamento de Tecnología de la Informática y Computación, Universidad de Alicante,
Apdo. 99. 03080 Alicante, Spain

Abstract. Recovering the shape of a class of objects requires establishing correct correspondences between manually or automatically annotated landmark points. In this study, we utilise a novel approach to automatically recover the shape of hand outlines from a series of 2D training images. Automated landmark extraction is accomplished through the use of the self-organising model the growing neural gas (GNG) network which is able to learn and preserve the topological relations of a given set of input patterns without requiring a priori knowledge of the structure of the input space. To measure the quality of the mapping throughout the adaptation process we use the topographic product. Results are given for the training set of hand outlines.

1

Towards Optimal Training of Cascaded Detectors

S. Charles Brubaker, Matthew D. Mullin, and James M. Rehg

College of Computing and GVV Center,

Georgia Institute of Technology, Atlanta, GA 30332

{brubaker, mdmullin, rehg}@cc.gatech.edu

Abstract. Cascades of boosted ensembles have become popular in the object detection community following their highly successful introduction in the face detector of Viola and Jones [1]. In this paper, we explore several aspects of this architecture that have not yet received adequate attention: decision points of cascade stages, faster ensemble learning, and stronger weak hypotheses. We present a novel strategy to determine the appropriate balance between false positive and detection rates in the individual stages of the cascade based on a probabilistic model of the overall cascade's performance. To improve the training time of individual stages, we explore the use of feature filtering before the application of Adaboost. Finally, we show that the use of stronger weak hypotheses based on CART can significantly improve upon the standard face detection results on the CMU-MIT data set.

1

Learning and Incorporating Top-Down Cues in Image Segmentation

Xuming He, Richard S. Zemel, and Debajyoti Ray

Department of Computer Science, University of Toronto

{hexm, zemel, debray}@cs.toronto.edu

Abstract. Bottom-up approaches, which rely mainly on continuity principles, are often insufficient to form accurate segments in natural images. In order to improve performance, recent methods have begun to incorporate top-down cues, or object information, into segmentation. In this paper, we propose an approach to utilizing category-based information in segmentation, through a formulation as an image labelling problem. Our approach exploits bottom-up image cues to create an over-segmented representation of an image. The segments are then merged by assigning labels that correspond to the object category. The model is trained on a database of images, and is designed to be modular: it learns a number of image contexts, which simplify training and extend the range of object classes and image database size that the system can handle. The learning method estimates model parameters by maximizing a lower bound of the data likelihood. We examine performance on three real-world image databases, and compare our system to a standard classifier and other conditional random field approaches, as well as a bottom-up segmentation method.

1

Learning to Detect Objects of Many Classes Using Binary Classifiers

Ramana Isukapalli¹, Ahmed Elgammal², and Russell Greiner³

1Lucent Technologies, Bell Labs Innovations, Whippany, NJ 07981, USA

2Rutgers University, New Brunswick, NJ 08854, USA

3University of Alberta, Edmonton, CA T6G 2E8, CA

Abstract. Viola and Jones [VJ] demonstrate that cascade classification methods can successfully detect objects belonging to a single class, such as faces. Detect-

ing and identifying objects that belong to any of a set of "classes", many class detection, is a much more challenging problem. We show that objects from each class can form a "cluster" in a "classifier space" and illustrate examples of such

clusters using images of real world objects. Our detection algorithm uses a "decision tree classifier" (whose internal nodes each correspond to a VJ classifier) to propose a class label for every sub-image of a test image (or reject it as a negative instance). If this reaches a leaf of this tree, we then pass through a

subsequent VJ cascade of classifiers, specific to the identified class, to determine

whether it is truly an instance of the proposed class. We perform several empirical studies to compare our system for detecting objects of any of M classes, to the obvious approach of running a set of M learned VJ cascade classifiers, one for each class of objects, on the same image. We found that the detection rates are comparable, and our many-class detection system is about as fast as running a single VJ cascade, and scales up well as the number of classes increases.

1

A Unifying Framework for Mutual Information

Methods for Use in Non-linear Optimisation

Nicholas Dowson and Richard Bowden

Centre for Vision Speed and Signal Processing,

University of Surrey, Guildford, GU2 7JW, UK

{n.dowson, r.bowden}@surrey.ac.uk

<http://www.ee.surrey.ac.uk/personal/n.dowson>

Abstract. Many variants of MI exist in the literature. These vary primarily in how the joint histogram is populated. This paper places the four main variants of MI: Standard sampling, Partial Volume Estimation (PVE), In-Parzen Windowing and Post-Parzen Windowing into a single mathematical framework. Jacobians and Hessians are derived in each case. A particular contribution is that the non-linearities implicit to standard sampling and post-Parzen windowing are explicitly dealt with. These non-linearities are a barrier to their use in optimisation. Side-by-side comparison of the MI variants is made using eight diverse data-sets, considering computational expense and convergence. In the experiments, PVE was generally the best performer, although standard sampling often performed nearly as well (if a higher sample rate was used). The widely used sum of squared differences metric performed as well as MI unless large occlusions and non-linear intensity relationships occurred. The binaries and scripts used for testing are available online.

1

Random Walks, Constrained Multiple Hypothesis

Testing and Image Enhancement

Noura Azzaboul^{1,2}, Nikos Paragios¹, and Frederic Guichard²

1M A S, Ecole Centrale de Paris,

Grande Voie des Vignes, Chatenay-Malabry, France

noura.azzabou@certis.enpc.fr, nikos.paragios@ecp.fr

<http://www.mas.ecp.fr>

2DxO Labs, 3, Rue Nationale, 92100 Boulogne, France

{nazzabou, fguichard}@dxo.com

<http://www.dxo.com>

Abstract. Image restoration is a keen problem of low level vision. In this paper

,
 we propose a novel - assumption-free on the noise model - technique based on random walks for image enhancement. Our method explores multiple neighbors sets (or hypotheses) that can be used for pixel denoising, through a particle filtering approach. This technique associates weights for each hypotheses according to its relevance and its contribution in the denoising process. Towards accounting for the image structure, we introduce perturbations based on local statistical properties of the image. In other words, particle evolution are controlled by the image structure leading to a filtering window adapted to the image content. Promising experimental results demonstrate the potential of such an approach.

1

From Tensor-Driven Diffusion to Anisotropic Wavelet Shrinkage

Martin Welk¹, Joachim Weickert¹, and Gabriele Steidl²

¹Mathematical Image Analysis Group,
 Faculty of Mathematics and Computer Science,
 Saarland University, 66041 Saarbrücken, Germany
 {welk, weickert}@mia.uni-saarland.de

<http://www.mia.uni-saarland.de>

²Faculty of Mathematics and Computer Science,
 A5 University of Mannheim, 68131 Mannheim, Germany
 steidl@math.uni-mannheim.de

<http://kiwi.math.uni-mannheim.de>

Abstract. Diffusion processes driven by anisotropic diffusion tensors are known to be well-suited for structure-preserving denoising. However, numerical implementations based on finite differences introduce unwanted blurring artifacts that deteriorate these favourable filtering properties. In this paper we introduce a novel

discretisation of a fairly general class of anisotropic diffusion processes on a 2-D grid. It leads to a locally semi-analytic scheme (LSAS) that is absolutely stable, simple to implement and offers an outstanding sharpness of filtered images. By showing that this scheme can be translated into a 2-D Haar wavelet shrinkage procedure, we establish a connection between tensor-driven diffusion and anisotropic wavelet shrinkage for the first time. This result leads to coupled

shrinkage rules that allow to perform highly anisotropic filtering even with the simplest wavelets.

1

SURF: Speeded Up Robust Features

Herbert Bay¹, Tinne Tuytelaars², and Luc Van Gool^{1,2}

¹ETH Zurich

{bay, vangool}@vision.ee.ethz.ch

²Katholieke Universiteit Leuven

{Tinne.Tuytelaars, Luc.Vangool}@esat.kuleuven.be

Abstract. In this paper, we present a novel scale- and rotation-invariant interest point detector and descriptor, coined SURF (Speeded Up Robust Features). It approximates or even outperforms previously proposed schemes with respect to repeatability, distinctiveness, and robustness, yet can be computed and compared much faster.

This is achieved by relying on integral images for image convolutions; by building on the strengths of the leading existing detectors and descriptors (in case, using a Hessian matrix-based measure for the detector, and a distribution-based descriptor); and by simplifying these methods to the essential. This leads to a combination of novel detection, description, and match

ing steps. The paper presents experimental results on a standard evaluation set, as well as on imagery obtained in the context of a real-life object recognition application. Both show SURF's strong performance.

1

Top-Points as Interest Points for Image Matching

B. Platel, E. Balmachnova, L.M.J. Florack[■], and B.M. ter Haar Romeny

Technische Universiteit Eindhoven, P.O. Box 513,

5600 MB Eindhoven, The Netherlands

{B.Platel, E.Balmachnova, L.M.J.Florack,

B.M.terHaarRomeny }@tue.nl

Abstract. We consider the use of top-points for object retrieval. These points are

based on scale-space and catastrophe theory, and are invariant under gray value scaling and offset as well as scale-Euclidean transformations. The differential properties and noise characteristics of these points are mathematically well understood. It is possible to retrieve the exact location of a top-point from any coarse

estimation through a closed-form vector equation which only depends on local derivatives in the estimated point. All these properties make top-points highly suitable as anchor points for invariant matching schemes. By means of a set of repeatability experiments and receiver-operator curves we demonstrate the performance of top-points and differential invariant features as image descriptors.

1

Machine Learning for High-Speed

Corner Detection

Edward Rosten and Tom Drummond

Department of Engineering,

Cambridge University, UK

{er258, twd20 }@cam.ac.uk

Abstract. Where feature points are used in real-time frame-rate applications, a high-speed feature detector is necessary. Feature detectors such as SIFT (DoG), Harris and SUSAN are good methods which yield high quality features, however they are too computationally intensive for use in real-time applications of any complexity. Here we show that machine learning can be used to derive a feature detector which can fully process live PAL video using less than 7% of the available processing time. By comparison neither the Harris detector (120%) nor the detection stage of SIFT (300%) can operate at full frame rate.

Clearly a high-speed detector is of limited use if the features produced are unsuitable for downstream processing. In particular, the same scene viewed from two different positions should yield features which correspond to the same real-world 3D locations[1]. Hence the second contribution of this paper is a comparison corner detectors based on this criterion applied to 3D scenes. This comparison supports a number of claims made elsewhere concerning existing corner detectors. Further, contrary to our initial expectations, we show that despite being principally constructed for speed, our detector significantly outperforms existing feature detectors according to this criterion.

1

Smooth Image Segmentation by Nonparametric

Bayesian Inference

Peter Orbanz and Joachim M. Buhmann

Institute of Computational Science, ETH Zurich

{porbanz, jbuhmann }@inf.ethz.ch

Abstract. A nonparametric Bayesian model for histogram clustering is proposed to automatically determine the number of segments when Markov Random Field constraints enforce smooth class assignments. Then nonparametric

ic nature of this model is implemented by a Dirichlet process prior to control the number of clusters. The resulting posterior can be sampled by a modification of a conjugate-case sampling algorithm for Dirichlet process mixture models. This sampling procedure estimates segmentations as efficiently as clustering procedures in the strictly conjugate case. The sampling algorithm can process both single-channel and multi-channel image data. Experimental results are presented for real-world synthetic aperture radar and magnetic resonance imaging data.

1

Shape Analysis and Fuzzy Control
for 3D Competitive Segmentation
of Brain Structures with Level Sets
Cybele Ciofalo and Christian Barillot
IRISA / CNRS, Team VisAGeS,
Campus de Beaulieu, 35042 Rennes Cedex, France
{Cybele.Ciofalo, Christian.Barillot}@irisa.fr
<http://www.iris.fr/visages/visages-eng.html>

Abstract. We propose a new method to segment 3D structures with competitive level sets driven by a shape model and fuzzy control. To this end, several contours evolve simultaneously toward previously defined targets. The main contribution of this paper is the original introduction of prior information provided by a shape model, which is used as an anatomical atlas, into a fuzzy decision system. The shape information is combined with the intensity distribution of the image and the relative position of the contours. This combination automatically determines the directional term of the evolution equation of each level set. This leads to a local expansion or contraction of the contours, in order to match the borders of their respective targets. The shape model is produced with a principal component analysis, and the resulting mean shape and variations are used to estimate the target location and the fuzzy states corresponding to the distance between the current contour and the target. By combining shape analysis and fuzzy control, we take advantage of both approaches to improve the level set segmentation process with prior information. Experiments are shown for the 3D segmentation of deep brain structures from MRI and a quantitative evaluation is performed on a 18 volumes dataset.

1

Variational Motion Segmentation with Level Sets
Thomas Brox¹, Andr es Bruhn², and Joachim Weickert²
¹CVPR Group, Department of Computer Science, University of Bonn,
R omerstr. 164, 53113 Bonn, Germany
brox@mia.uni-saarland.de
²Mathematical Image Analysis Group, Faculty of Mathematics and Computer Science,
Saarland University, Building 27, 66041 Saarbr ucken, Germany
{bruhn, weickert}@mia.uni-saarland.de

Abstract. We suggest a variational method for the joint estimation of optic flow and the segmentation of the image into regions of similar motion. It makes use of the level set framework following the idea of motion competition, which is extended to non-parametric motion. Moreover, we automatically determine an appropriate initialization and the number of regions by means of recursive two-phase splits with higher order region models. The method is further extended to the spatiotemporal setting and the use of additional cues like the gray value or color for the segmentation. It need not fear a quantitative comparison to pure optical flow estimation techniques: For the popular Yosemite sequence with clouds we obtain the

currently most accurate result. We further uncover a mistake in the ground truth. Coarsely correcting this, we get an average angular error below 1 degree.

1

Ellipse Fitting with Hyperaccuracy

Kenichi Kanatani

Department of Computer Science,

Okayama University, Okayama 700-8530, Japan

kanatani@suri.it.okayama-u.ac.jp

Abstract. For fitting an ellipse to a point sequence, ML (maximum likelihood) has been regarded as having the highest accuracy. In this paper, we demonstrate the existence of a "hyperaccurate" method which outperforms ML. This is made possible by error analysis of ML followed by subtraction of high-order bias terms. Since ML nearly achieves the theoretical accuracy bound (the KCR lower bound), the resulting improvement is very small. Nevertheless, our analysis has theoretical significance, illuminating the relationship between ML and the KCR lowerbound.

1

A Physically-Motivated Deformable Model

Based on Fluid Dynamics

Andrei C. Jalba and Jos B.T.M. Roerdink

Institute for Mathematics and Computing Science,

University of Groningen, P.O. Box 800,

9700 AV, Groningen, The Netherlands

{andrei, roe}@cs.rug.nl

Abstract. A novel deformable model for image segmentation and shape recovery is presented. The model is inspired by fluid dynamics and is based on a flooding simulation similar to the watershed paradigm. Unlike most watershed methods, our model has a continuous formulation, being described by two partial differential equations. In this model, different fluids, added by placing density (dye) sources manually or automatically, are attracted towards the contours of the objects of interest by an image force. In contrast to the watershed method, when different fluids meet they may mix. When the topographical relief of the image is flooded, the interfaces separating homogeneous fluid regions can be traced to yield the object contours. We demonstrate the flexibility and potential of our model in two experimental settings: shape recovery using manual initializations and automated segmentation.

1

Video and Image Bayesian Demosaicing

with a Two Color Image Prior

Eric P. Bennett¹, Matthew Uyttendaele², C. Lawrence Zitnick²,

Richard Szeliski², and Sing Bing Kang²

¹The University of North Carolina at Chapel Hill, Chapel Hill, NC

²Microsoft Research, Redmond, WA

Abstract. The demosaicing process converts single-CCD color representations of one color channel per pixel into full per-pixel RGB. We introduce a Bayesian technique for demosaicing Bayer color filter array patterns that is based on a statistically-obtained two color per-pixel image prior. By modeling all local color behavior as a linear combination of two fully specified RGB triples, we avoid color fringing artifacts while preserving sharp edges. Our grid-less, floating-point pixel location architecture can process both single images and multiple images from video within the same framework, with multiple images providing denser color samples and therefore better color reproduction with reduced aliasing. An initial clustering is performed to determine the underlying local two color model surrounding each pixel. Using a product of Gaussians statistical model, the underlying linear blending ratio of the two representative

colors at each pixel is estimated, while simultaneously providing noise reduction. Finally, we show that by sampling the image model at a finer resolution than the source images during reconstruction, our continuous demosaicing technique can super-resolve in a single step.

1

Generalized Multi-sensor Planning

Anurag Mittal

Dept of Computer Science and Engg.,
Indian Institute of Technology Madras,
Chennai-600036, India

Abstract. Vision systems for various tasks are increasingly being deployed. Although significant effort has gone into improving the algorithms for such tasks, there has been relatively little work on determining optimal sensor configurations.

This paper addresses this need. We specifically address and enhance the state-of-the-art in the analysis of scenarios where there are dynamically occurring objects

capable of occluding each other. The visibility constraints for such scenarios are

analyzed in a multi-camera setting. Also analyzed are other static constraints such

as image resolution and field-of-view, and algorithmic requirements such as stereo

reconstruction, face detection and background appearance. Theoretical analysis with the proper integration of such visibility and static constraints leads to a generic

framework for sensor planning, which can then be customized for a particular task.

Our analysis can be applied to a variety of applications, especially those involving

randomly occurring objects, and include surveillance and industrial automation. Several examples illustrate the wide applicability of the approach.

1

Variational Shape and Reflectance Estimation

Under Changing Light and Viewpoints

Neil Birkbeck¹, Dana Cobzas¹, Peter Sturm², and Martin Jagersand¹

¹Computer Science, University of Alberta, Canada

{birkbeck, dana, jag}@cs.ualberta.ca

²INRIA Rhone-Alpes, France

peter.sturm@inrialpes.fr

Abstract. Fitting parameterized 3D shape and general reflectance models to 2D image data is challenging due to the high dimensional nature of the problem. The proposed method combines the capabilities of classical and photometric stereo, allowing for accurate reconstruction of both textured and non-textured surfaces. In particular, we present a variational method implemented as a PDE-driven surface evolution interleaved with reflectance estimation. The surface is represented on an adaptive mesh allowing topological change. To provide the input data, we have designed a pipeline that simultaneously acquires both viewpoint and light variation while minimizing self-shadowing. Our capture method is feasible for real-world application as it requires a moderate amount of input data and processing time. In experiments, models of people and everyday objects were captured from a few dozen images taken with a consumer digital camera. The captured process recovers a photo-consistent model of spatially varying Lambertian and specular reflectance and a highly accurate geometry.

1

Specularity Removal in Images and Videos:

A PDE Approach

Satya P. Mallick¹, Todd Zickler², Peter N. Belhumeur³, and David J. Kriegman¹

¹Computer Science and Engineering, University of California at San Diego, CA 92093

²Engineering and Applied Sciences, Harvard University, Cambridge, MA 02138

³Computer Science, Columbia University, New York, NY 10027

Abstract. We present a unified framework for separating specular and diffuse reflection components in images and videos of textured scenes. This can be used for specularity removal and for independently processing, filtering, and recombining the two components. Beginning with a partial separation provided by an illumination-dependent color space, the challenge is to complete the separation using spatio-temporal information. This is accomplished by evolving a partial differential equation (PDE) that iteratively erodes the specular component at each pixel. A family of PDEs appropriate for differing image sources (still images vs. videos), differing prior information (e.g., highly vs. lightly textured scenes), or differing prior computations (e.g., optical flow) is introduced. In contrast to many other methods, explicit segmentation and/or manual intervention are not required. We present results on high-quality images and video acquired in the laboratory in addition to images taken from the Internet. Results on the latter demonstrate robustness to low dynamic range, JPEG artifacts, and lack of knowledge of illuminant color. Empirical comparison to physical removal of specularities using polarization is provided. Finally, an application termed dichromatic editing is presented in which the diffuse and the specular components are processed independently to produce a variety of visual effects.

1

Carved Visual Hulls for Image-Based Modeling

Yasutaka Furukawaland Jean Ponce^{1,2}

¹Department of Computer Science, University of Illinois at Urbana Champaign, USA

²Département d'Informatique, Ecole Normale Supérieure, Paris, France

{yfurukaw, ponce}@cs.uiuc.edu

Abstract. This article presents a novel method for acquiring high-quality solid models of complex 3D shapes from multiple calibrated photographs. After the purely geometric constraints associated with the silhouettes found in each image have been used to construct a coarse surface approximation in the form of a visual hull, photoconsistency constraints are enforced in three consecutive steps: (1) the rims where the surface grazes the visual hull are first identified through dynamic programming; (2) with the rims now fixed, the visual hull is carved using graph cuts to globally optimize the photoconsistency of the surface and recover its main features; (3) an iterative (local) refinement step is finally used to recover fine surface details. The proposed approach has been implemented, and experiments with six real data sets are presented, along with qualitative comparisons with several state-of-the-art image-based-modeling algorithms.

1

What Is the Range of Surface Reconstructions from a Gradient Field?

Amit Agrawal¹, Ramesh Raskar², and Rama Chellappa¹

¹Center for Automation Research, University of Maryland, College Park, MD, USA 20742

{aagrawal, rama }@cfar.umd.edu
2Mitsubishi Electric Research Labs (MERL),
201 Broadway, Cambridge, MA, USA 02139
raskar@merl.com

Abstract. We propose a generalized equation to represent a continuum of surface reconstruction solutions of a given non-integrable gradient field. We show that common approaches such as Poisson solver and Frankot-Chellappa algorithm are special cases of this generalized equation. For a $N \times N$ pixel grid, the subspace of all integrable gradient fields is of dimension $N^2 - 1$. Our framework can be applied to derive a range of meaningful surface reconstructions from this high dimensional space. The key observation is that the range of solutions is related to the degree of anisotropy in applying weights to the gradients in the integration process. While common approaches use isotropic weights, we show that by using a progression of spatially varying anisotropic weights, we can achieve significant improvement in reconstructions. We propose (a) α -surfaces using binary weights, where the parameter α allows trade off between smoothness and robustness, (b) M-estimators and edge preserving regularization using continuous weights and (c) Diffusion using affine transformation of gradients. We provide results on photometric stereo, compare with previous approaches and show that anisotropic treatment discounts noise while recovering salient features in reconstructions.

1

Practical Global Optimization for
Multiview Geometry

Sameer Agarwal¹, Manmohan Krishna Chandraker¹, Fredrik Kahl²,
David Kriegman¹, and Serge Belongie¹

¹University of California, San Diego, CA 92093, USA

{sagarwal, mkchandraker, kriegman, sjb }@cs.ucsd.edu

²Lund University, Lund, Sweden

fredrik@maths.lth.se

Abstract. This paper presents a practical method for finding the provably global optimal solution to numerous problems in geometry including multiview triangulation, camera resection and homography estimation. Unlike traditional methods which may get trapped in local minima due to the non-convex nature of these problems, this approach provides a theoretical guarantee of global optimality. The formulation relies on recent developments in fractional programming and the theory of convex underestimation as well as a unified framework for minimizing the standard L_2 -norm of reprojection errors which is optimal under Gaussian noise as well as the more robust L_1 -norm which is less sensitive to outliers. The efficacy of our algorithm is empirically demonstrated by good performance on experiments for both synthetic and real data. An open source MATLAB toolbox that implements the algorithm is available to facilitate further research.

1

Perspective n-View Multibody

Structure-and-Motion Through Model Selection

Konrad Schindler, James U, and Hanzi Wang

Institute for Vision Systems Engineering,

Monash University, Clayton, 3800 VIC, Australia

{Konrad.Schindler, James.U, Hanzi.Wang }@eng.monash.edu.au

Abstract. Multi-body structure-and-motion (MSaM) is the problem to establish the multiple-view geometry of an image sequence of a 3D scene, where the scene consists of multiple rigid objects moving relative to each

other. Solutions have been proposed for several restricted settings, such as only two views, a single projection, and perspective projection of linearly moving points. We give a solution for sequences of several images, full perspective projection, and general rigid motion. It can deal with the fact that there is feedback of responses changes over time, and is robust to outliers. The proposed solution is based on Monte-Carlo sampling and clustering of two-view motions, linking them through the sequence, and model selection to yield the best explanation for the entire sequence.

1

Confocal Stereo

Dept. of Computer Science, University of Toronto

Abstract. We present confocal stereo, a new method for computing 3D

shape by controlling the focus using a depth sensor. The method is

specifically designed for reconstructing scenes with high geometric com-

plexity or fine-scale texture. To achieve this, we introduce the confocal

constancy property, which states that as the sensor varies, the

pixel intensity of a visible in-focus scene point will vary in a scene-

independent way, that can be predicted by prior radiometric calibration.

The only requirement is that incoming radiance within the cone

subtended by the largest aperture is nearly constant. First, we develop

a detailed lens model that factors out the distortions in high resolution

SLR cameras (12MP or more) with large-aperture lenses (e.g., f1.2).

This allows us to assemble an $A \times F$ aperture-focus image (AFI) for

each pixel, that collects the undistorted measurements over all apertures

and focus settings. In the AFI representation, confocal constancy

reduces to color comparisons within regions of the AFI, a natural set of

metrics that can be evaluated separately for each pixel. We propose

two such metrics and present initial reconstruction results for complex

scenes.

1
